

Deep Learning

Alexander Panin, Andrey Ustyuzhanin, Alena Zarodnyuk

2021, LaSalle University



LAMBDA • HSE

Recap of day3

KNN classification

Random Forest

Boosting, Bagging

ROC curve, AUC

X-validation

Regularization

Decision Tree

Overfitting

Loss function

Purity of a leaf

Accuracy/Recall/Precision

Confusion Matrix

RMSE

Linear Regression

Model: $x \longrightarrow wx + b \longrightarrow y^{\text{pred}}$

Objective function:
$$L = \frac{1}{N} \sum_i (y_i - y_i^{\text{pred}})^2$$

Optimization (exact):
$$w = (X^T \cdot X)^{-1} X^T y$$


Linear Regression

Model: $x \longrightarrow wx + b \longrightarrow y^{\text{pred}}$

Objective function:
$$L = \frac{1}{N} \sum_i (y_i - y_i^{\text{pred}})^2$$

Optimization (iterative):
$$w_{i+1} := w_i - \alpha \cdot \frac{\partial L}{\partial w_i}$$

Logistic Regression




$$p(y|x) = \sigma(Wx + b) = \frac{1}{1 + e^{-(Wx+b)}}$$

Objective function ?

Logistic Regression

Model:



Objective function:

$$L = -\frac{1}{N} \sum_i y_i \cdot \log p(y|x_i) + (1 - y_i) \cdot \log(1 - p(y|x_i))$$


Optimization(iterative): **same as linear regression**

Recap: Gradient Descent


Update:

$$w_{i+1} := w_i - \alpha \cdot \frac{\partial L}{\partial w_i}$$


- α – learning rate $\alpha \ll 1$
- L – loss function



Nonlinear dependencies



What we have




What we want

How to learn that?

Feature extraction

Loss: same as linear/logistic regression

Model:



Training:

$$\underset{\theta_1}{\operatorname{argmin}} L$$


What if...



Features would tune to your problem automatically!

What do we want, exactly?

Model:




Objective:

?

$$\underset{\theta_1}{\operatorname{argmin}} L$$

What do we want, exactly?

Model:




Joint training:

$$\underset{\theta_1, \theta_2}{\operatorname{argmin}} \quad L$$

Okay, how do we extract features?

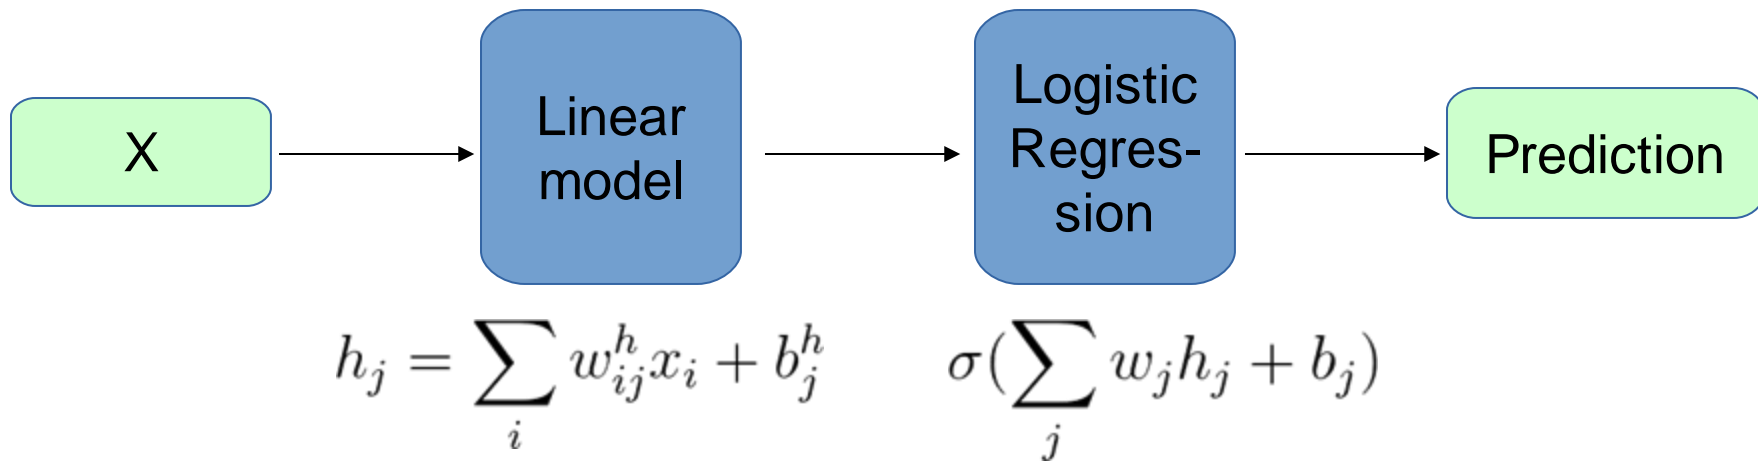
Try linear

Model:




Try linear

Model:



Try linear

Model:



$$h_j = \sum_i w_{ij}^h x_i + b_j^h$$

$$\sigma\left(\sum_j w_j h_j + b_j\right)$$

Output:
$$p(y|x) = \sigma\left(\sum_j w_j \left(\sum_i w_i^h x_i + b_i^h\right) + b_j\right)$$

Is it any better than logistic regression?

Try linear

$$p(y|x) = \sigma\left(\sum_j w_j \left(\sum_i w_i^h x_i + b_i^h\right) + b_j\right)$$


$$p(y|x) = \sigma\left(\sum_i \left[\sum_j w_j w_i^h\right] x_i + \sum_i b_i^h + b_j\right)$$

$$\hat{w}_i = \left[\sum_j w_j w_i^h\right] x_i; \quad \hat{b} = \sum_i b_i^h + b_j$$

$$p(y|x) = \sigma\left(\sum_i \hat{w}_i x_i + \hat{b}_i\right)$$

Add nonlinearity



Model:



$$p(y|x) = \sigma\left(\sum_j w_j \sigma\left(\sum_i w_i^h x_i + b_i^h\right) + b_j\right)$$

Add nonlinearity


Model:



Add nonlinearity

- $f(a) = 1/(1+e^a)$
- $f(a) = \tanh(a)$


- $f(a) = \max(0, a)$
- $f(a) = \log(1+e^a)$



a

Training the monster

Model:




Output:
$$p(y|x) = \sigma\left(\sum_j w_j \sigma\left(\sum_i w_i^h x_i + b_i^h\right) + b_j\right)$$

Training:

?!

Training the monster

Model:



Output:
$$p(y|x) = \sigma\left(\sum_j w_j \sigma\left(\sum_i w_i^h x_i + b_i^h\right) + b_j\right)$$

Training: gradient descent!
$$w_{i+1} := w_i - \alpha \cdot \frac{\partial L}{\partial w_i}$$

Backpropagation

- **TL;DR:** backprop = chain rule*

$$\frac{\partial f(g(x))}{\partial x} = \frac{\partial f(g(x))}{\partial g(x)} \cdot \frac{\partial g(x)}{\partial x}$$

Backpropagation

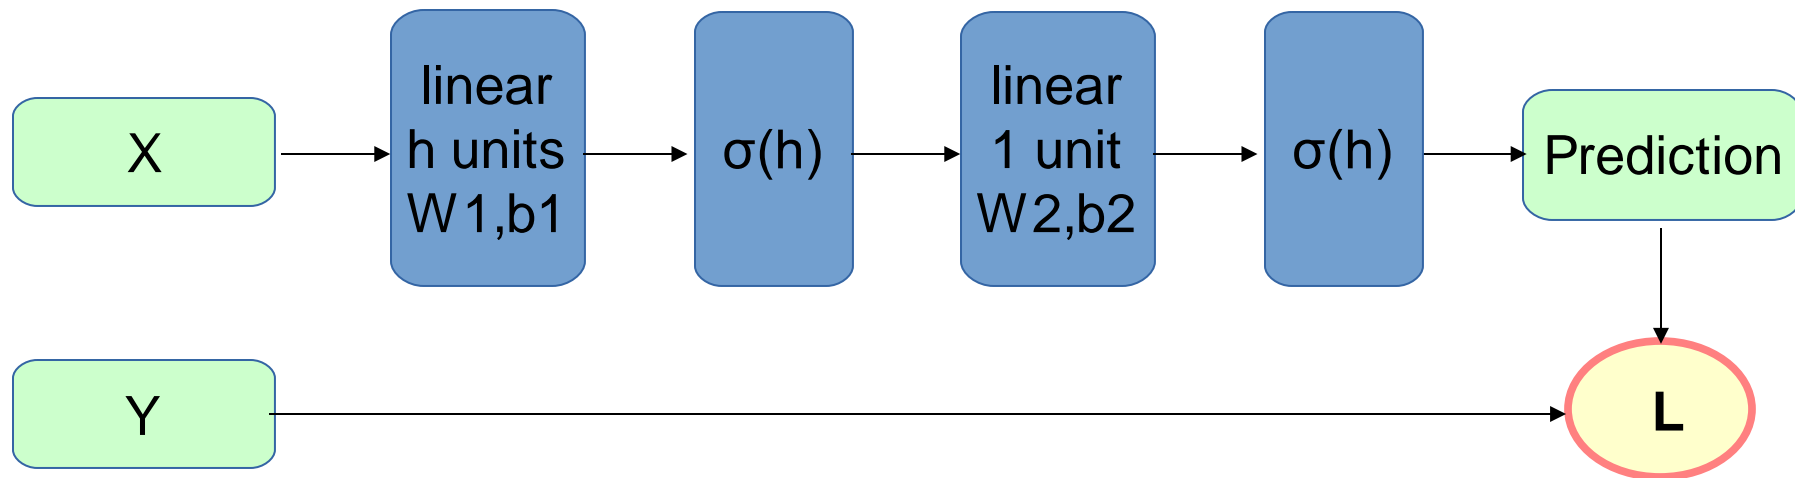
- **TL;DR:** backprop = chain rule*

$$\frac{\partial f(g(x))}{\partial x} = \frac{\partial f(g(x))}{\partial g(x)} \cdot \frac{\partial g(x)}{\partial x}$$

* g and x can be vectors/vectors/tensors

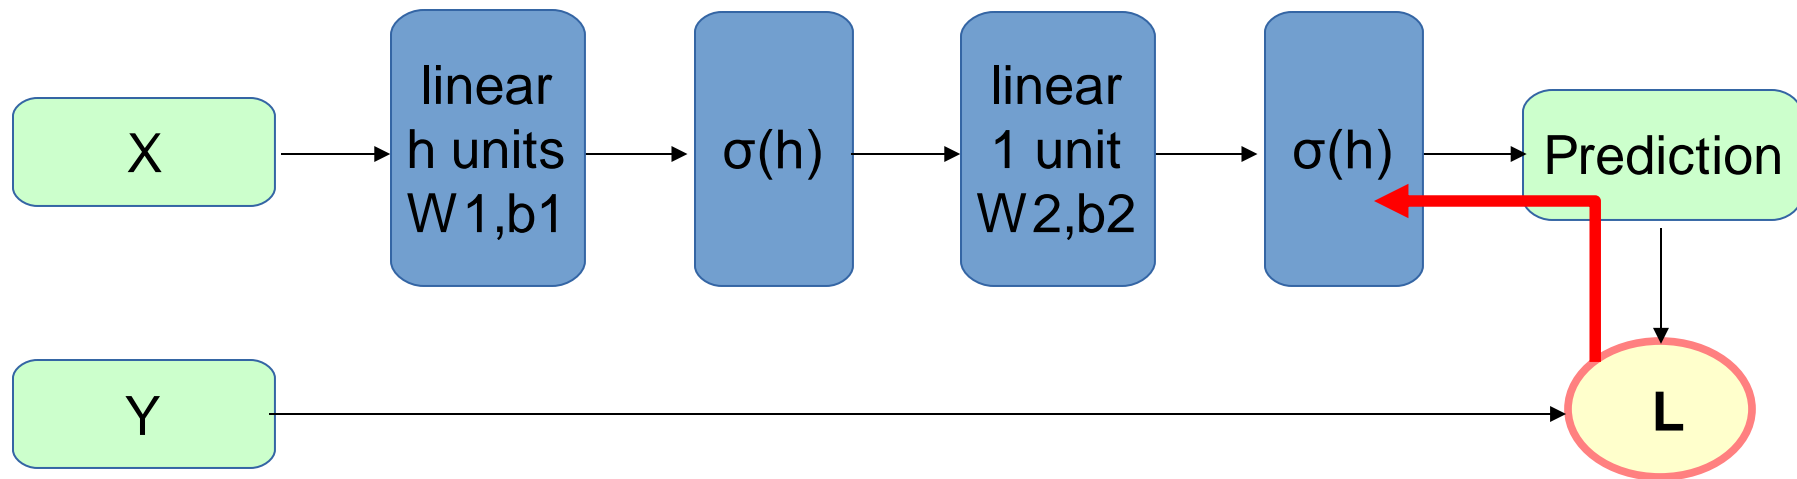


Backpropagation



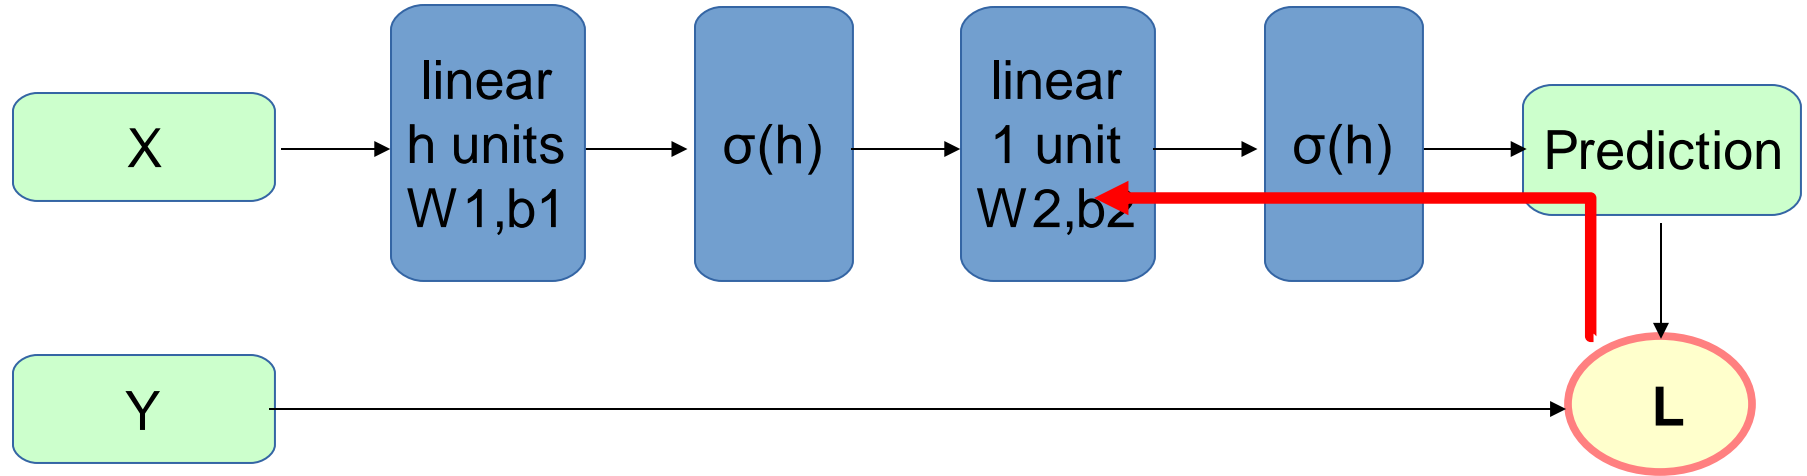
$$\frac{\partial L(\text{linear}_2(\sigma(\text{linear}_1(x))))}{\partial W_1} = \dots$$

Backpropagation



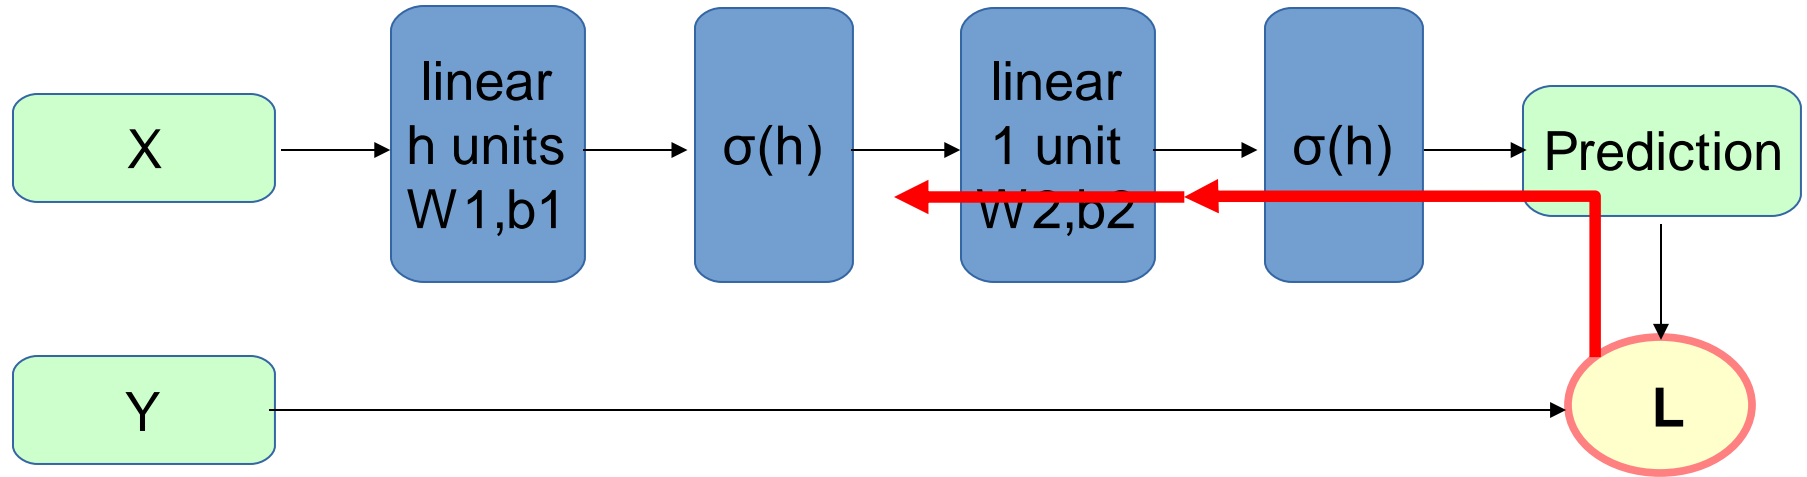
$$\frac{\partial L}{\partial W_1} = \frac{\partial L}{\partial \sigma_2} .$$

Backpropagation



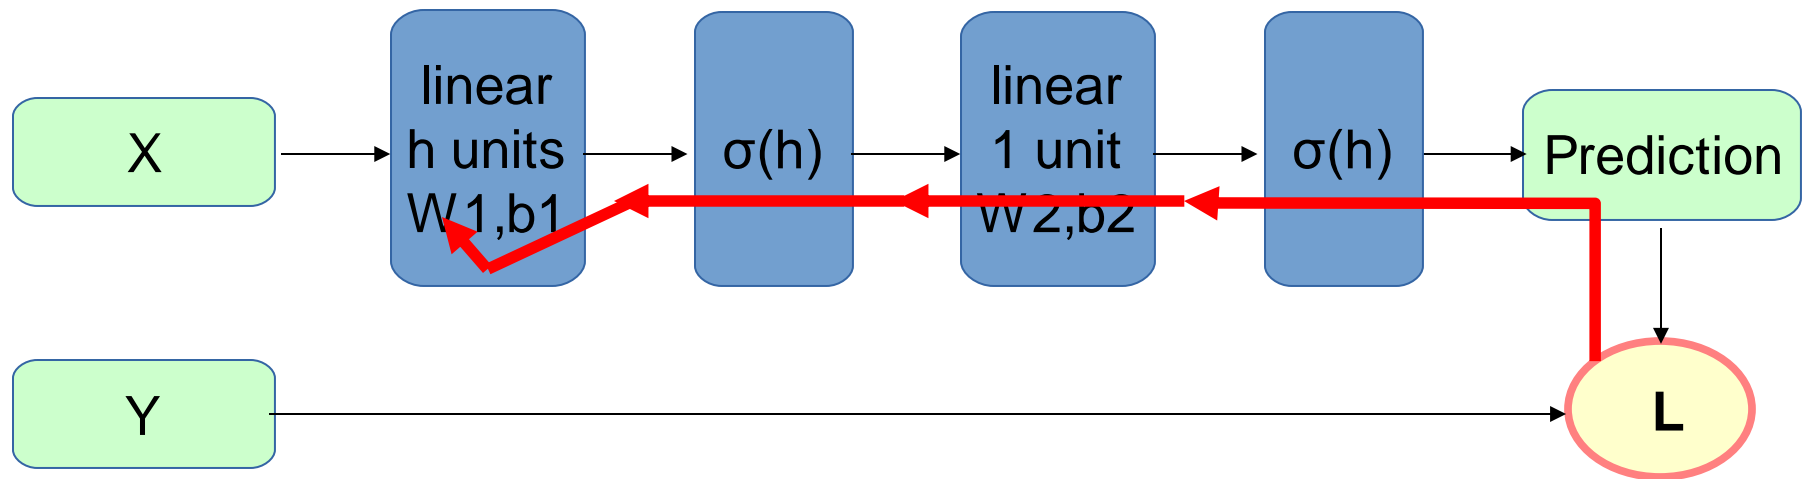
$$\frac{\partial L}{\partial W_1} = \frac{\partial L}{\partial \sigma_2} \cdot \frac{\partial \sigma_2}{\partial linear_2} \cdot$$

Backpropagation



$$\frac{\partial L}{\partial W_1} = \frac{\partial L}{\partial \sigma_2} \cdot \frac{\partial \sigma_2}{\partial linear_2} \cdot \frac{\partial linear_2}{\partial \sigma_1} \cdot \dots$$

Backpropagation



$$\frac{\partial L}{\partial W_1} = \frac{\partial L}{\partial \sigma_2} \cdot \frac{\partial \sigma_2}{\partial linear_2} \cdot \frac{\partial linear_2}{\partial \sigma_1} \cdot \frac{\partial \sigma_1}{\partial linear_1} \cdot \frac{\partial linear_1}{\partial W_1}$$

Matrix derivatives

Let's compute:

$$\frac{\partial L(X \times W + b)}{\partial X} = \frac{\partial L(X \times W + b)}{\partial [X \times W + b]} \times \boxed{\text{What?}}$$

Variable shapes:

 X

[batch size, features]

 W

[features, outputs]

 b

[outputs]

$$\frac{\partial L(X \times W + b)}{\partial X}$$

[batch size, features]

$$\frac{\partial L(X \times W + b)}{\partial [X \times W + b]}$$

[batch size, outputs]

Matrix derivatives

Let's compute:

*Hint: 1. figure out scalar case,
2. match shapes for matrices*

$$\frac{\partial L(X \times W + b)}{\partial X} = \frac{\partial L(X \times W + b)}{\partial [X \times W + b]} \times \boxed{\text{What?}}$$

Variable shapes:

X

[batch size, features]

$\frac{\partial L(X \times W + b)}{\partial X}$

[batch size, features]

W

[features, outputs]

$\frac{\partial L(X \times W + b)}{\partial [X \times W + b]}$

[batch size, outputs]

b

[outputs]

Matrix derivatives

Let's compute:

$$\frac{\partial L(X \times W + b)}{\partial X} = \frac{\partial L(X \times W + b)}{\partial [X \times W + b]} \times W^T$$

Variable shapes:

X

[batch size, features]

$\frac{\partial L(X \times W + b)}{\partial X}$

[batch size, features]

W

[features, outputs]

$\frac{\partial L(X \times W + b)}{\partial [X \times W + b]}$

[batch size, outputs]

b

[outputs]

Matrix derivatives

Let's compute:

$$\frac{\partial L(X \times W + b)}{\partial W} =$$

What?

Variable shapes:

X

[batch size, features]

$$\frac{\partial L(X \times W + b)}{\partial X}$$

[batch size, features]

W

[features, outputs]

$$\frac{\partial L(X \times W + b)}{\partial [X \times W + b]}$$

[batch size, outputs]

b

[outputs]

Matrix derivatives

Let's compute:

$$\frac{\partial L(X \times W + b)}{\partial W} = X^T \times \frac{\partial L(X \times W + b)}{\partial [X \times W + b]}$$

Variable shapes:

 X

[batch size, features]

 $\frac{\partial L(X \times W + b)}{\partial X}$ ∂X

[batch size, features]

 W

[features, outputs]

 $\frac{\partial L(X \times W + b)}{\partial [X \times W + b]}$ $\partial [X \times W + b]$

[batch size, outputs]

 b

[outputs]

Cheat sheet for seminar

Gradient of $\sum_i \log p(y_i|x_i, w) = \sum_i \text{gradient} \log p(y_i|x_i, w)$

linear over X : $\frac{\partial L}{\partial [X \times W + b]} \times W^T$


linear over W : $\frac{1}{\|X\|} \cdot X^T \times \frac{\partial L}{\partial [X \times W + b]}$

sigmoid : $\frac{\partial L}{\partial \sigma(x)} \cdot [\sigma(x) \cdot (1 - \sigma(x))]$


Works for any kind of x
(scalar, vector, matrix, tensor)

Not actual neurons :)

- Neurons output in “spikes”,
not real numbers
- No one knows for sure
how they “train”
- There are other cell types
e.g. neuroglial cells




Neuroglial Cells of the CNS



Connectionist phrasebook


- Layer – a building block for NNs :
 - “Dense layer”: $f(x) = Wx+b$
 - “Nonlinearity layer”: $f(x) = \sigma(x)$
 - Input layer, output layer
 - A few more we gonna cover later
- Activation – layer output
 - i.e. some intermediate signal in the NN
- Backpropagation – a fancy word for “chain rule”

Connectionist phrasebook



“Train it via backprop!”

Connectionist phrasebook



How do we train it?


Potential caveats?

Problems with deep learning

- Hardcore overfitting
- No “golden standard” for architecture
- Computationally heavy

Back to neural networks

Model:



Training:




see the [demo](#)

Gradient Descent

Update:

$$w_{i+1} := w_i - \alpha \cdot \frac{\partial L}{\partial w_i}$$

- α – learning rate $\alpha \ll 1$
- L – loss function




Can we do better?

Newton-Raphson

Parameter update

$$w_{i+1} = w_i - \alpha \cdot H^{-1} \frac{\partial L}{\partial w_i}$$

Hessian: $\mathbf{H} = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}.$



Green: Gradient Descent

Red: Newton-Raphson


Any drawbacks?

Newton-Raphson

Parameter update

$$w_{i+1} = w_i - \alpha \cdot H^{-1} \frac{\partial L}{\partial w_i}$$

Hessian: $\mathbf{H} = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}.$



Green: Gradient Descent

Red: Newton-Raphson


**Inverting \mathbf{H} might
be infeasible**

Stochastic gradient descent

Approximate objective with samples

$$L = \frac{1}{N} \sum_i f(x_i, y_i, w) = E_{i \sim U(1, N)} f(x_i, y_i, w)$$

$$L \approx f(x_i, y_i, w) |_{i \sim U(1, N)}$$



Stochastic gradient descent


Approximate objective with samples

$$L = \frac{1}{N} \sum_i f(x_i, y_i, w) = E_{i \sim U(1, N)} f(x_i, y_i, w)$$

$$L \approx f(x_i, y_i, w) |_{i \sim U(1, N)}$$

Update:

$$w_{i+1} = w_i - \alpha \cdot \frac{\partial \hat{L}}{\partial w} \quad \text{where } \hat{L} = f(x_i, y_i, w) |_{i \sim U(1, N)}$$



SGD with momentum

Idea: average gradient over consecutive steps
aka “add inertia”

$$w_0 := 0; \nu_0 := 0$$

$$\nu_{i+1} := \alpha \cdot \frac{\partial L}{\partial w} + \mu \cdot \nu_i$$

$$w_{i+1} := w_i - \nu_{i+1}$$

SGD with momentum

Idea: average gradient over consecutive steps
aka “add inertia”

$$w_0 := 0; \nu_0 := 0$$



$$\nu_{i+1} := \alpha \cdot \frac{\partial L}{\partial w} + \mu \cdot \nu_i \quad // \text{ reuse gradient from previous steps}$$

$$w_{i+1} := w_i - \nu_{i+1}$$

SGD with momentum

Idea: average gradient over consecutive steps
(see *this [demo](#)*)

$$\nu_{i+1} := \alpha \cdot \frac{\partial L}{\partial w} + \mu \cdot \nu_i \qquad w_{i+1} := w_i - \nu_{i+1}$$



RMSProp

Idea: adapt learning rate to the slope of objective function
large gradient = go slower, small gradient = go faster

$$w_0 := 0; ms_0 := 0$$

$$ms_{i+1} := \gamma \cdot ms_i + (1 - \gamma) \cdot \left\| \frac{\partial L}{\partial w} \right\|^2$$

$$w_{i+1} := w_i - \frac{\alpha}{\sqrt{ms_{i+1} + \epsilon}} \frac{\partial L}{\partial w}$$

RMSProp

Idea: adapt learning rate to the slope of objective function
large gradient = go slower, small gradient = go faster

$$w_0 := 0; ms_0 := 0$$

moving average norm²

$$ms_{i+1} := \gamma \cdot ms_i + (1 - \gamma) \cdot \left\| \frac{\partial L}{\partial w} \right\|^2$$

$$w_{i+1} := w_i - \frac{\alpha}{\sqrt{ms_{i+1} + \epsilon}} \frac{\partial L}{\partial w}$$

Adam

Idea: combine momentum and rmsprop in one method
the “default” optimizer, see the [paper](#) for details

$$w_0 := 0; ms_0 := 0$$

moving average norm²


$$ms_{i+1} := \gamma \cdot ms_i + (1 - \gamma) \cdot \left\| \frac{\partial L}{\partial w} \right\|^2$$

$$w_{i+1} := w_i - \frac{\alpha}{\sqrt{ms_{i+1} + \epsilon}} \frac{\partial L}{\partial w}$$

TL;DR stochastic optimization

Tips & tricks

- Adam works fine out of the box
- One can usually beat adam with tuned sgd+momentum
- Tuning may take long
- Everyone has his favorite optimizer :)



Stuff we won't cover


First order

- . Adam flavors: NAdam, Adamax, QHAdam
- . Adagrad, Radagrad, Adadelata - alternative adaptive lr

Approx second order

- . BFGS, L-BFGS
- . K-FAC


Initialization, symmetry problem



Initialize with zeros
 $W := 0$

What will the first
step look like?

Initialization, symmetry problem



- Break the symmetry!
- Initialize with random numbers!

$W := N(0, 0.01)?$
 $W := U(-0.1, 0.1)?$


[read more](#)

Nuff said

Let's go implement that!




And now let's differentiate




- near state-of-the-art in image classification
- 5+ types of layers
- parallel branches with independent losses
- few hundred megabytes of weights

And now let's differentiate



- near state-of-the-art in image classification
- 5+ types of layers
- parallel branches with independent losses
- few hundred megabytes of weights

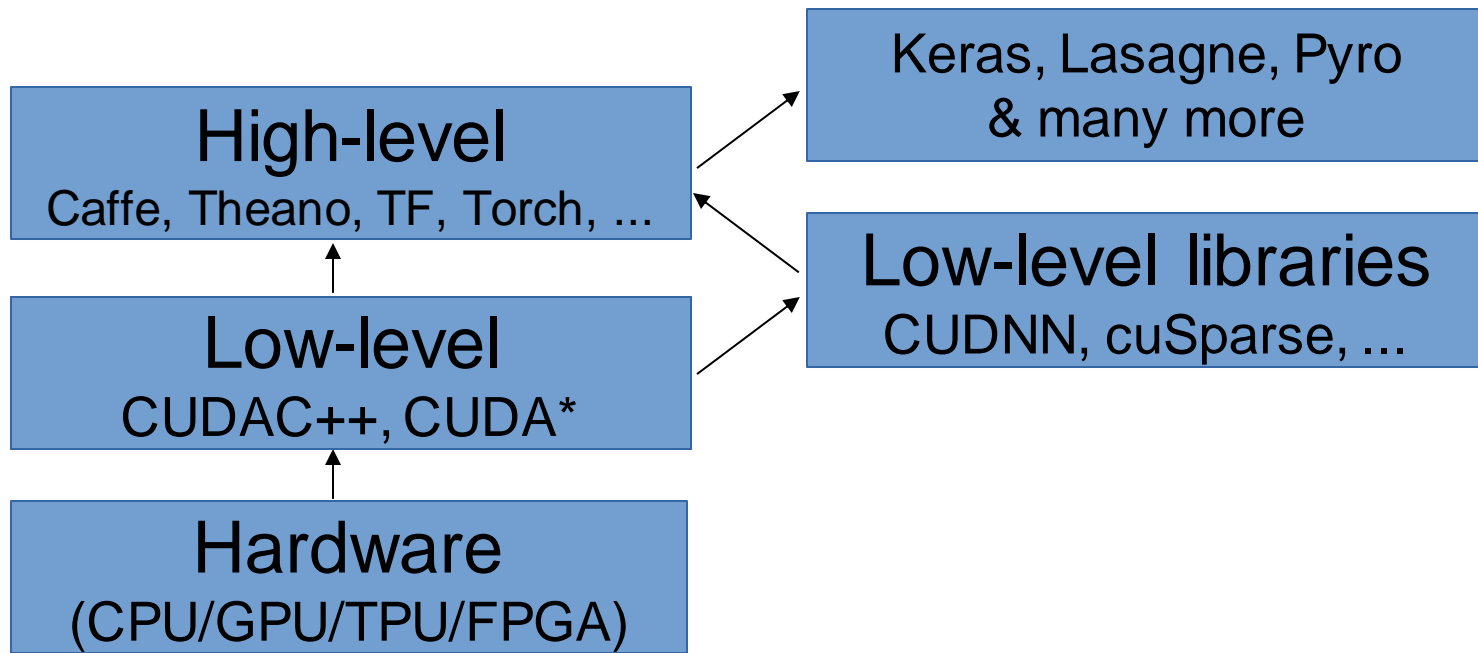
Deep learning frameworks



- near state-of-the-art in image classification
- 5+ types of layers
- parallel branches with independent losses
- few hundred megabytes of weights

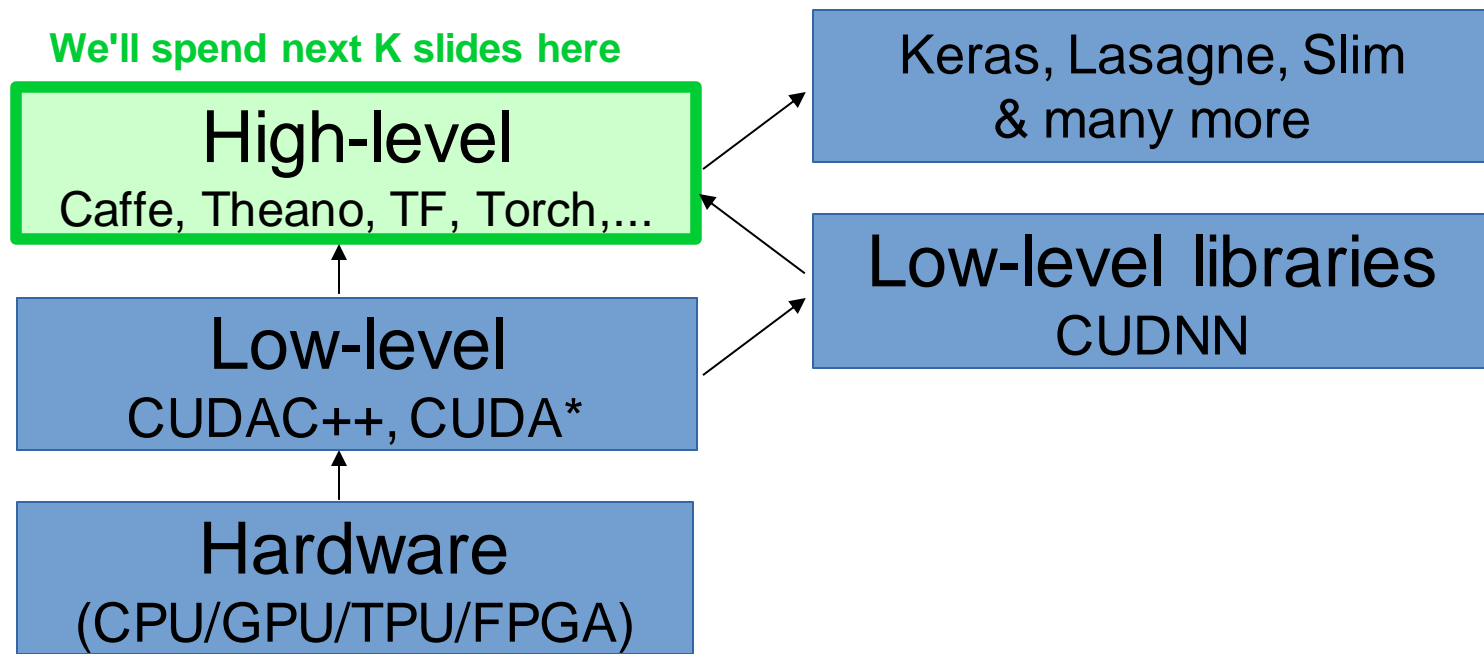
Deep learning frameworks

- Core idea: helps you define and train neural nets



Deep learning frameworks

- Core idea: helps you define and train neural nets




Deep learning frameworks

- Layer-based frameworks:
 - Same idea as in our hand-made neural net

Deep learning frameworks

- Layer-based frameworks:
 - Same idea as in our hand-made neural net
 - this one - <http://bit.ly/2w9kAHm>



Deep learning frameworks

Caffe

```
name: "LeNet"
layer {
  name: "conv1"
  type: "Convolution"
  bottom: "data"
  top: "conv1"
  param {lr_mult: 1}
  param {lr_mult: 2}
  convolution_param {
    num_output: 20
    kernel_size: 5
    stride: 1
    weight_filler {
      type: "xavier"
    }
  }
}
```

[130 lines](#)

....

You define model in config file
by stacking layers (left)

Then run train script:

```
caffe train -solver
examples/mnist/lenet_solver.proto
txt
```

Deep learning frameworks

Caffe

```
name: "LeNet"
layer {
  name: "conv1"
  type: "Convolution"
  bottom: "data"
  top: "conv1"
  param {lr_mult: 1}
  param {lr_mult: 2}
  convolution_param {
    num_output: 20
    kernel_size: 5
    stride: 1
    weight_filler {
      type: "xavier"
    }
  }
}
```

[130 lines](#)

- + Easy to deploy (C++)
- + A lot of pre-trained models (model zoo)
- Model as protobuf
- Hard to build new layers
- Hard to debug


Still used for
computer vision

Computation graphs

What does your CPU do
when you write this?

```
a = x * y  
b = a + z  
c = np.sum(b)
```

Computation graphs




`a = x * y`

`b = a + z`

`c = np.sum(b)`

Idea: let's define
this graph explicitly!

Computation graphs



- + Automatic gradients!
- + Easy to build new layers
- + We can optimize the Graph
- Graph is static during training
- Need time to compile/optimize
- Hard to debug

Dynamic graphs

- Chainer, DyNet, Pytorch

A graph is created on the fly

```
from torch.autograd import Variable

x = Variable(torch.randn(1, 10))
prev_h = Variable(torch.randn(1, 20))
W_h = Variable(torch.randn(20, 20))
W_x = Variable(torch.randn(20, 10))
```



Dynamic graphs

- Chainer, DyNet, Pytorch



- + Can change graph on the fly
- + Can get value of any tensor at any time
(easy debugging)
- Hard to optimize graphs
(especially large graphs)
- Still early development

Researchers love these!

Dynamic graphs



Andrej Karpathy ✓

@karpathy

Following



I've been using PyTorch a few months now and I've never felt better. I have more energy. My skin is clearer. My eye sight has improved.

Researchers love them!